




Pandas

- **Pandas** is a python package containing data structures to make working with labeled data intuitive and easy
 - The two primary data structures are *Series* and *DataFrames*
- The **Pandas** package also contains several useful functions for
 - **Plotting data**
 - **Manipulating data structures**
 - **Data manipulation**



Series

- **Series** is one of the primary data structures offered in Pandas
 - They are a one dimensional labeled array
 - They can hold any data type (e.x. string, int, python objects, etc)
- They can be created from *array-like objects, iterable objects, dictionaries, scalar values*

```
s = pd.Series(  
    {'Corn Flakes': 100.0,  
     'Almond Delight': 110.0,  
     'Cinnamon Toast Crunch': 120.0,  
     'Cocoa Puff': 110.0}    )
```

s

```
-----  
Corn Flakes          100.0  
Almond Delight       110.0  
Cinnamon Toast Crunch 120.0  
Cocoa Puff           110.0  
dtype: float64
```

Constructor

```
pandas.Series(data=None, index=None, dtype=None, name=None)
```

Parameter	Expected value
data	array-like, Iterable, dict, or scalar value
index	array-like or Index (1d)
dtype	str, numpy.dtype, or ExtensionDtype,
name	str

Available dtypes

Pandas dtype	Python Type	NumPy type
Object	Str or mixed values	String_, unicode_, mixed types
int64	int	int_, int8, int16, int32, int64, uint8, uint16, uint32, uint64
float64	float	float_, float16, float32, float64
bool	bool	bool_
datetime64	N/A	datetime64[ns]
timedelta[ns]	N/A	N/A
category	N/A	N/A

Series Attributes

→ Series.dtype

- Returns the overall data type of the underlying object

→ Series.size

- Returns the number of elements in the underlying data

→ Series.index

- Return an array containing the index

```
import pandas as pd

s = pd.Series(
    {'Corn Flakes': 100.0,
     'Almond Delight': 110.0,
     'Cinnamon Toast Crunch': 120.0,
     'Cocoa Puff': 110.0} )

print(s.dtype)
print(s.size)
print(s.index)
-----
..float64
..4
..array(['Corn Flakes', 'Almond Delight', 'Cinnamon
Toast Crunch', 'Cocoa Puffs'])
```

Series Attributes

→ Series.hasnans

- Returns True if there are nans in the series

→ Series.empty

- Returns true if the Series is empty

→ Series.name

- Returns the name of the Series (the column name)

```
import pandas as pd

nans_series = pd.Series([1,2,3,None,5])
empty_series = pd.Series()
named_series = pd.Series(['Apple', 'Pear',
                           'Fig'], name= 'Fruits')

print(nans_series.hasnans)
print(empty_series.empty)
print(named_series.name)
-----
True
True
'Fruits'
```

Conversions

- **Series.to_numpy()**
 - Will return an array containing the Series data
- **Series.astype(dtype, copy=True, errors='raise')**
 - Returns a Series with data type cast to specified dtype
- **Series.copy()**
 - Returns a deep copy of a Series

```
import pandas as pd

num_series = pd.Series([1,2])
numpy_arr = num_series.to_numpy()
copy_series = num_series.to_numpy()
print(type(num_series))
print(num_series.astype(float))
print(num_series.hasnans)
-----
numpy.ndarray
0      1.0
1      2.0
```


Viewing Series

→ **Series.iloc[<index>]**

- Returns data at the specified index position

→ **Series.loc[<label>]**

- Returns data at the specified label. Uses index if there's no label
- Is also aliased to [**<label>**]

→ The data in a pandas series is mutable

```
import pandas as pd

supermarket_prices = pd.Series(
    {'Garlic': 0.50,
     'Eggs': 7.50,
     'Ramen': 1.00,
     'Cabbage': 2.00})

print(supermarket_prices.iloc[1])
print(supermarket_prices.loc['Ramen'])
print(supermarket_prices['Cabbage'])

-----
0.50
1.00
2.00
```

Viewing Series

- **Series.head(n)**
 - Shows the first n rows
- **Series.tail(n)**
 - Shows the last n rows
- **Series[series == condition]**
 - We can also search a series based on a condition
 - Returns series whose elements fulfill condition

```
import pandas as pd

programming_languages = pd.Series(['Python',
                                   'Java', 'C++', 'JavaScript', 'React', 'R'], 1)

print('Head: ')
print(programming_languages.head(1))
print('Tail: ')
print(programming_languages.tail(2))
```

Head:

0	Python
---	--------

Tail:

4	React
5	R

Math

- **Series.count()**
 - Returns a count of non NA values
- **Series.max()**
 - Returns the largest value
- **Series.mean()**
 - Returns the average of the series
- **Series.min()**
 - Returns the minimum value

```
import pandas as pd

grocery= pd.Series(
    [5.50, 2.30, 6.20, 3.20] index=['Cheese',
    'Potato', 'Butter', 'Cream'])

print(grocery.count())
print(grocery.max())
print(grocery.mean())
print(grocery.min())

-----
4
6.2
4.3
2.3
```

Math

- **Series.std()**
 - Returns the mean of the series
- **Series.describe()**
 - Returns the max, min, count, mean and std in a series
- **Series.abs()**
 - Returns the absolute value of a series

```
print(grocery.std())  
print(grocery.describe())
```

```
-----  
1.849324200890693
```

```
count    4.000000  
mean     4.300000  
std      1.849324  
min      2.300000  
25%      2.975000  
50%      4.350000  
75%      5.675000  
max      6.200000  
dtype: float64
```

Math

→ **Series.mode()**

- Returns a Series containing the mode(s)

→ **Series.prod()**

- Returns the product of the Series

→ **Series.value_counts()**

- Returns a Series containing a count of each unique value

```
import pandas as pd

tips= pd.Series([300, 250, 250, 225, 75, 250])
print(tips.mode())
print(tips.prod())
print(tips.value_counts())
-----
0      250
79101562500000
250      3
300      1
225      1
75       1
dtype: int64
```

Manipulation

- **Series.drop(labels=None, inplace=False)**
 - Returns a Series with specified index labels removed
- **Series.drop_duplicates(keep='first', inplace=False)**
 - Returns Series with duplicate values removed
 - Choices for keep parameter are first, last, or False (removes all duplicates)

```
programming = pd.Series(['Python', 'Java',  
                        'C++', 'Java'])  
no_c = programming_languages.drop(2)  
no_java = programming_languages.drop_duplicates(keep=  
False)  
print(no_react)  
print(no_java)  
-----  
0    Python  
1     Java  
3     Java  
dtype: object  
0    Python  
2     C++  
dtype: object
```

Manipulation

→ **Series.isna()**

- Returns a Series of booleans mapping if a value is missing

→ **Series.dropna(inplace=False)**

- Returns a Series with all missing values removed

→ **Series.fillna(value=None, inplace=False)**

- Returns a Series with all missing values replaced with 'value'

```
none_series= pd.Series([1, 2, None])  
print(none_series.isna())  
print(none_series.isna())  
print(none_series.fillna(100))
```

```
-----  
0    False  
1    False  
2     True  
dtype: bool  
0     1.0  
1     2.0  
dtype: float64  
0     1.0  
1     2.0  
2    100.0  
dtype: float64
```

Manipulation

→ **Series.groupby(by=None, dropna=True)**

- If Series is multi Indexed, specify the index in by, otherwise we pass series name
- Returns groupby object that is waiting for a math function

→ **Series.unique()**

- Returns a numpy array of unique values

```
import pandas as pd

programming = pd.Series(['Python', 'Java',
                        'C++', 'Java'])
print(programming.groupby(programming))
print(programming.groupby(programming).count())
print(programming.unique())

-----
<pandas.core.groupby.generic.SeriesGroupBy object at
0x7f3cd4372370>
C++      1
Java     2
Python   1
dtype: int64
array(['Python', 'Java', 'C++'], dtype=object)
```


Manipulation

→ **Series.map(arguments)**

- Used to substitute each value in a Series with another value, that is derived from a function, dictionary, or Series.
- A passed function will be used to edit each value
- A passed dictionary or Series replaces values based on the key value pairs

```
import pandas as pd

s= pd.Series(['Cat', 'dog',
             np.nan, 'rabbit'])
print(s)
print(s.map('I am a {}'.format()) )
```

```
0      cat
1      dog
2      NaN
3    rabbit
dtype: object
0      I am a cat
1      I am a dog
2      I am a nan
3      I am a rabbit
dtype: object
```

Manipulation

- **Pandas.concat(objs*, ignore_index=False)**
- Combines 2 or more Pandas objects by stacking them
 - Pandas objects are joined in the order passed
 - If ignore_index is set to True, a new index is created

```
import pandas as pd

playable_char= pd.Series(['Mario','Luigi','Yoshi'] )
enemies= pd.Series(['Bowser','Waluigi','Boo'])
characters=pd.concat([enemies,playable_char],ignore_index=True)
print(characters)

-----
0      Bowser
1    Waluigi
2        Boo
3      Mario
4      Luigi
5      Yoshi
dtype: object
```

Series I/O functions for Series

Series Function	Function Definition
<code>Series.to_csv(path=None)</code>	Write object to a csv file.
<code>Series.to_json(path=None)</code>	Convert the object to a JSON string.
<code>Series.to_pickle(path)</code>	Compresses a file and writes it to memory
<code>Series.to_list()</code>	Returns a list of the values
<code>Series.to_dict()</code>	Convert Series to {label -> value} dict or dict-like object.

Student Exercise

- Create the following Pandas Series
 - ◆ Employee **ID**: type `int`
 - ◆ Employee **Name**: type `str`
 - ◆ Employee **Salary**: type `int`
 - ◆ Employee **Department**: type `str`
 - ◆ Employee **Start Date**: type `datetime`
 - ◆ **Currently Employed**: type `bool`
- Populate each Series with an equal number of entries
- Set the Index of the other series to be the Employee id



DataFrame



DataFrame

- **DataFrame** is the other primary data structures offered in Pandas
 - They are 2 dimensional tables
 - Each column in a DataFrame is a Series
- They can be created from *array-like objects*, *iterable objects*, *dictionaries*, or *DataFrame*
- Many Series functions will also work on DataFrame objects

```
Characters = pd.DataFrame(  
    {'Name': ['Luffy', 'Sanji', 'Chopper'],  
    'Age': [17, 19, 8],  
    'Role': ['Captain', 'Cook', 'Doctor']} )
```

Characters

	Name	Age	Role
0	Luffy	17	Captain
1	Sanji	19	Cook
2	Chopper	8	Doctor

Constructor

```
pd.DataFrame(data=None, index=None, columns=None, dtype=None)
```

Parameter	Expected value
data	Array-like, iterable, dict, or a DataFrame. Dictionary keys are column names
index	array-like or Index (1d)
columns	Labels to use for column names
dtype	Data type to enforce. Can only choose 1 dtype

Attributes

- Just like in Series, we can see different DataFrame attributes including
 - `DataFrame.size`
 - `DataFrame.index`
 - `DataFrame.empty`
- We can also get the data types of each column using
 - `DataFrame.dtypes`

```
Characters = pd.DataFrame(  
    {'Name': ['Luffy', 'Sanji', 'Chopper'],  
    'Age': [17, 19, 8],  
    'Role': ['Captain', 'Cook', 'Doctor']} )  
  
print(Characters.size)  
print(Characters.dtypes)  
-----  
9  
Name      object  
Age       float64  
Role      object  
dtype: object
```


Attributes

- We can get the the shape, or how many rows and columns are in a dataframe with
 - **DataFrame.shape**
- Shape is returned in tuple with rows listed first and columns listed second

```
Characters = pd.DataFrame(  
    {'Name': ['Luffy', 'Sanji', 'Chopper'],  
    'Role': ['Captain', 'Cook', 'Doctor']})
```

```
print(Characters)  
print(Characters.shape)
```

```
-----  
   Name  Role  
0  Luffy  Captain  
1  Sanji   Cook  
2  Chopper Doctor
```

```
(3,2)
```

Indexing

- To select a column as a whole, use the command
 - `df[<column label>]`
 - We can also use `df.column`
- To select particular rows by their labels, use
 - `df.loc[<label>]`
 - The labels passed can be a single label, a list of labels, a slice object with label, or a boolean

```
Characters = pd.DataFrame(  
    {'Name': ['Luffy', 'Sanji', 'Chopper'], 'Age':  
    [17, 19, 8], 'Role': ['Captain', 'Cook', 'Doctor']} )  
Characters = Characters.set_index('Name')  
print(Characters['Role'])  
print(Characters.loc[2])  
-----  
0    Captain  
1      Cook  
2     Doctor  
Name: Role, dtype: object  
Name    Chopper  
Age      8  
Role     Doctor  
Name: 2, dtype: object
```

Indexing

- To select particular rows by their index position, use
 - `df.iloc[<index>]`
 - Can be passed a integer, list of integers, slice object with integers, or boolean
- Both loc and iloc will raise errors if a label or index is not found

```
Characters = pd.DataFrame(  
    {'Name': ['Luffy', 'Sanji', 'Chopper'],  
    'Age': [17, 19, 8],  
    'Role': ['Captain', 'Cook', 'Doctor']} )  
Characters = Characters.set_index('Name')  
print(Characters.loc['Sanji'])  
print(Characters.iloc[1])
```

```
-----  
Age      19  
Role     Cook  
Name: Sanji, dtype: object  
Age      19  
Role     Cook  
Name: Sanji, dtype: object
```

Indexing

- We can also use `DataFrame.query` to index particular rows of data
 - `DataFrame.query(expression, inplace=False)`
- You can refer to variables in the environment by prefixing them with an '@' character like `@a + b`

```
Characters = pd.DataFrame(  
    {'Name': ['Luffy', 'Sanji', 'Chopper'],  
    'Age': [17, 19, 8],  
    'Role': ['Captain', 'Cook', 'Doctor']} )  
num = 10  
print(Characters.query('Age > @num'))
```

```
-----  
   Name  Age  Role  
0  Luffy   17  Captain  
1  Sanji   19   Cook
```

Conversions

- We can also call and apply these functions across DataFrames
 - `DataFrame.to_numpy()`
 - `DataFrame.copy()`
- We can change each individual column's datatype by passing a dictionary to
 - `DataFrame.astype({column_name: data type})`

```
import pandas as pd

num_series = pd.DataFrame([1,2])
numpy_arr = num_series.to_numpy()
copy_series = num_series.copy()
print(type(numpy_arr))
print(num_series.astype(float))
print(copy_series is num_series )

-----
<class 'numpy.ndarray'>
0
0 1.0
1 2.0
False
```

Student Exercise

- Use the **Employee Series** you created to make a **DataFrame**
- This **DataFrame** should have the **Employee ID** as its index, and the other **Series** columns
- Try different methods of creating your **DataFrame**!
 - ◆ Appending **Series** to an existing DF
 - ◆ Creating from a **list of Series**
 - ◆ Creating from a **list of dictionaries**
 - ◆ Creating from a **dictionary of Series**
- Access data from your **DataFrame** using the accessor methods



Math

- All Series math functions can also be applied across **DataFrame** as a whole
- Both DataFrame and Series also support binary operators

```
nums = pd.DataFrame({'A': [1, 5, 20]})  
print(nums.sum())  
print(nums+1)
```

```
-----  
      A  
0      2  
1      6  
2     21  
  
A      26  
dtype: int64
```

Changing Data

→ We can change each individual element of a DataFrame with

- **DataFrame.applymap(Func)**
- The passed function must be callable
- The function will be applied to element
- A **DataFrame** is returned

```
import pandas as pd

num_series = pd.DataFrame([1,2.12],
                           [3.356,4.567])

print(num_series.applymap(lambda x:
                           len(str(x))))
```

	0
3.356	3
4.567	4

Changing Data

- We can group data exactly like how we group data in SQL
 - `DataFrame.groupby(by= None)`
- The groupby function returns a `DataFrameGroupBy` object which contains information about the groups
- Must use some math or aggregate function to view data as DataFrame

```
df = pd.DataFrame({'Animal': ['Falcon',  
                              'Falcon', 'Parrot', 'Parrot'],  
                  'Max_Speed':  
                  [380,3700,24,26]} )
```

```
print(df.groupby('Animal').mean())
```

```
-----  
                Max Speed  
Animal  
Falcon          375.0  
Parrot           25.0
```

Changing Data

We can also use an aggregate function in order to use one or more operations across a specified axis or grouped data

- `DataFrame.agg(function=None, axis=0)`
- `DataFrame.GroupBy.agg(function=None)`
- `.agg` accepts functions, function names, a list of functions/function names, or a dictionary of axis labels and function names

```
df = pd.DataFrame({'Animal': ['Falcon',  
                              'Falcon', 'Parrot', 'Parrot'],  
                  'Max_Speed':  
                  [380,3700,24,26]})  
  
print(df.groupby('Animal').agg(['sum', 'max'])  
)
```

```
-----  
              Max Speed  
              sum  max  
Animal  
Falcon          750  380  
Parrot           50   26
```

Combining Data

→ We can combine the columns of 2 DataFrame using the following

- `DataFrame.join(other, on=None, how='left', lsuffix="", rsuffix=")`
- This acts similar to SQL join, joining the columns of one dataframe with another dataframe on the index

```
ind = ['r1', 'r2', 'r3', 'r4']
df = pd.DataFrame({'Courses': ['Spark', 'PySpark', 'Python', 'Pandas'], 'Fee': [20000, 25000, 22000, 30000]}, index = ind )

ind2 = ['r1', 'r6', 'r3', 'r5']
df2 = pd.DataFrame({'Courses': ['Spark', 'Java', 'Python', 'Go'], 'Discount': [2000, 2300, 1200, 2000]}, index = ind2 )
print(df.join(df2, lsuffix= '_tech1', rsuffix= '_tech2', how = 'inner')
```

```
-----
      Courses_tech1  Fee  Courses_tech2  Discount
r1      Spark      20000      Spark      2000
r3      Python      22000      Python      1200
```

Combining Data

- We can pass a list of DataFrames
- When we join 2 DataFrames, use the parameters *'lsuffix'* and *'rsuffix'* to mark which **DataFrame** columns originally came from
- We can also specify the type of join we'd like including *'left'*, *'right'*, *'outer'*, and *'inner'*, but it will default to *inner*

```
ind = ['r1', 'r2', 'r3', 'r4']
df = pd.DataFrame({'Courses': ['Spark', 'PySpark', 'Python', 'Pandas'], 'Fee': [20000, 25000, 22000, 30000]}, index = ind )

ind2 = ['r1', 'r6', 'r3', 'r5']
df2 = pd.DataFrame({'Courses': ['Spark', 'Java', 'Python', 'Go'], 'Discount': [2000, 2300, 1200, 2000]}, index = ind2 )
print(df.join(df2, lsuffix= '_tech1', rsuffix= '_tech2', how = 'inner')
```

```
-----
      Courses_tech1  Fee  Courses_tech2  Discount
r1      Spark      20000      Spark      2000
r3      Python      22000      Python      1200
```

Changing Data

- **DataFrame.merge(right, how='inner', on=None, left_on=None, right_on=None, left_index=False, right_index=False, suffixes=('_x', '_y'))**
- Lets us join 2 DataFrames together on non index columns

```
ind = ['r1', 'r2', 'r3', 'r4']
df = pd.DataFrame({'Courses': ['Spark', 'PySpark', 'Python', 'Pandas'], 'Fee': [20000, 25000, 22000, 30000]}, index = ind )

ind2 = ['r1', 'r6', 'r3', 'r5']
df2 = pd.DataFrame({'Courses': ['Spark', 'Java', 'Python', 'Go'], 'Discount': [2000, 2300, 1200, 2000]}, index = ind2 )
print(df.merge(df2, on= 'Courses')
```

	Courses	Fee	Discount
0	Spark	20000	2000
1	Python	22000	1200

Changing Data

→ right

- This is where we specify the dataframe on the right of the join

→ how='inner',

- Here we can specify the type of join between 'inner', 'outer', 'left', 'right', 'cross'

```
ind = ['r1', 'r2', 'r3', 'r4']
df = pd.DataFrame({'Courses': ['Spark', 'PySpark', 'Python', 'Pandas'], 'Fee': [20000, 25000, 22000, 30000]}, index = ind )

ind2 = ['r1', 'r6', 'r3', 'r5']
df2 = pd.DataFrame({'Courses': ['Spark', 'Java', 'Python', 'Go'], 'Discount': [2000, 2300, 1200, 2000]}, index = ind2 )
print(df.merge(df2, on= 'Courses')
```

```
-----
   Courses  Fee  Discount
0   Spark  20000      2000
1  Python  22000      1200
```

Changing Data

→ *on=None*,

- Specify the column to join on. This must be in both DataFrames
- If blank, defaults to the intersection of the columns

→ *left_on=None, right_on=None*

- Specify the column to join on for each left and right DataFrame

```
ind = ['r1', 'r2', 'r3', 'r4']
df = pd.DataFrame({'Courses': ['Spark', 'PySpark', 'Python', 'Pandas'], 'Fee': [20000, 25000, 22000, 30000]}, index = ind )

ind2 = ['r1', 'r6', 'r3', 'r5']
df2 = pd.DataFrame({'Courses': ['Spark', 'Java', 'Python', 'Go'], 'Discount': [2000, 2300, 1200, 2000]}, index = ind2 )
print(df.merge(df2, on= 'Courses')
```

	Courses	Fee	Discount
0	Spark	20000	2000
1	Python	22000	1200

Changing Data

→ *left_index=False,*
right_index=False

- A bool value for if the join should use the index instead of a column

→ *suffixes=('_x', '_y')*

- A sequence of 2 suffixes to add to the columns from the left and right DataFrame respectively

```
ind = ['r1', 'r2', 'r3', 'r4']
df = pd.DataFrame({'Courses': ['Spark', 'PySpark', 'Python', 'Pandas'], 'Fee': [20000, 25000, 22000, 30000]}, index = ind )

ind2 = ['r1', 'r6', 'r3', 'r5']
df2 = pd.DataFrame({'Courses': ['Spark', 'Java', 'Python', 'Go'], 'Discount': [2000, 2300, 1200, 2000]}, index = ind2 )
print(df.merge(df2, on= 'Courses')
```

	Courses	Fee	Discount
0	Spark	20000	2000
1	Python	22000	1200

Student Exercise

- Find the **sum, average, minimum, and maximum** salaries.
 - ◆ Then find these measurements by department
- Create a **DataFrame** of only active employees
- Construct a separate **DataFrame** of **Departments** and merge it onto your Employees **DataFrame**
- Everyone gets a raise! Increase the salaries of all employees by 10%





FIN